

Fuzzing zkVMs

A report on progress toward secure scaling of Ethereum

Cody Gunton - Ethereum Foundation

Outline

- The basic challenges of fuzzing
- SNARK pipelines and where we might fuzz them
- First focus: program execution and ISA compliance
- Findings for security and observations on tooling
- Future directions

Fuzzing?

The art of throwing random inputs at a program to try to break it.

Factors for success:


- Quality of generated inputs
 - Do the inputs "go deep"?
 - What is the "throughput" of novel inputs?
- Speed of input generation

Fuzzing is not formal verification

Formal verification : "prove" that your code conforms to a spec.

- A major ongoing initiative <https://verified-zkevm.org/>
- Formal verification is not foolproof
- High degree of assurance has a high startup cost
 - Want tools that are easier to implement
 - Want to find and fix bugs before formal verification

Generic zkVM pipeline



RISC-V Primer

RISC-V is a family of ISAs $rv\{XLEN\}\{exts\}$ where

- $XLEN$ is 32 or 64 (the register size)
- $exts$ is a list of letters describing the available structures
 - Popular in SNARK settings: $rv32im$
 - Ideal for Go, Java,... compiled to SNARKs: $rv64gc$
 - Want this for client diversity!

First Investigation: Execution

How can we gain confidence that a RISC-V implementation is correct?

Naive differential fuzzing:

- Generate a random RISC-V program
- Execute on a trusted reference (Sail model), recording checkpoints
- Execute on target, recording the same
- Compare

Testing

- Unit tests: <https://github.com/riscv-software-src/riscv-tests>
 - Issue #368: doesn't work directly for rv32im
 - The testing framework assume old version of base set (rv32i)
- Architecture tests <https://github.com/riscv-non-isa/riscv-arch-test>
 - Reasonable framework for differential fuzzing
 - Implemented for some; results mostly good; should keep looking
 - Found a circuit bug


Test Binary Generation

- Generate then compile (e.g. CSmith for C programs)
- Generate RISC-V directly with risc-dv
 - Assumes you have SystemVerilog and want to simulate
 - Issue 979: not compatible with Verilator (\approx only free option)
- Cascade: looks good; claims to solve throughput issue well

Smarter Input Generation

- Described approach is robust but coarse
 - How do you know what you're testing?
- Introspective approaches more work but very powerful
 - Coverage guidance: track stats about which lines were executed
 - Need to code template for fuzzer to find interesting cases
 - Can use existing tools (e.g. Honggfuzz); but FFIs?

The Pipeline Again



Circuit Correctness

- A huge concern!
- Security firms have many sophisticated tools in development.
- We care about completeness bugs because we care about liveness
(need to be able to prove valid transactions are valid)
- Fuzzing can also help us find soundness bugs (cf Hochrainer,
Isychev, Wüstholtz, Christakis "Fuzzy Processing Pipelines...")

Prover Correctness

- Can focus on commonly used libraries
- Can extend circuit generation pipeline
 - Will be slow though
- Can also focus on components (sumcheck? PCS?)?

Verifier Correctness

- Seems easier to fuzz the verifier
 - Semantically simpler space of inputs
 - Weaker performance requirements \Rightarrow simpler tech stack (?)
 - Verification is very fast \Rightarrow can cover many inputs

Thanks for your attention!

References

- Solt, M., Lees, P., Koens, T., & Gerlitz, O., *Cascade: CPU Fuzzing via Intricate Program Generation*, USENIX Security '24, 2024.
- Hochrainer, C., Isychev, A., Wüstholtz, V., & Christakis, M., *Fuzzing Processing Pipelines for Zero-Knowledge Circuits*, arXiv:2411.02077 (Nov 2024).